

1 Introduction

This report seeks to validate the statement by Ahuja et al. (1993, p. 143) as referenced in the assignment guide. It details two primary tasks undertaken for this purpose:

- The development of Julia functions to implement two variations of the Bellman-Ford algorithm. The first variation employs a First-In-First-Out (FIFO) queue, while the second uses a double-ended queue (deque) as its data structure.
- The execution of experiments on networks with varying structures to evaluate the performance of these algorithms.

No AI tools were utilized for the development and experimental analysis. However, AI assistance was sought for refining the report's language and expression, specifically employing GPT-4.0 for sentence structure and expression enhancement and Grammarly for grammar checks.

Two Julia scripts are submitted:

- **ShortestPathProblems.jl**: Key functions are defined in this script.
- **main.jl**: This script is the main file performing the experiment by importing functions defined in 'ShortestPathProblems.jl.' It generates instances, solves them, and saves the results. The generated instances are saved in the 'instances' folder as .csv files, and the results are saved in the 'results' folder as .csv files.

The remainder of this report is organized into four sections:

- Section 2 introduces the codes implementing the algorithms, providing a brief overview.
- Section 3 provides the experimental design and the generation of instances briefly, describing the methodology and parameters used for testing the algorithms.
- Section 4 presents the experiments' results, analyzing the algorithms' performance across various testing cases.
- Section 5 concludes the report by summarizing the findings and their implications.

2 Algorithms & codes summary

The function implementing the algorithms is in 'ShortestPathProblems.jl' script:

```

176 function ModifiedBellmanFordAlgorithm(A, s::Int64 = 1, use_dequeue::Bool = true)
177     ## preprocessing
178     calV = maximum(A[:, 1:2])
179     calE = size(A)[1]
180     pnt = [findfirst(x->x==v, A[:, 1]) for v in 1:1:(calV+1)]
181     pnt[calV+1] = calE+1
182     tpv1 = findall(x->isnothing(x), pnt)
183     tpv2 = deleteat!(collect(1:1:(calV+1)), tpv1)
184     for v in tpv1
185         pnt[v] = pnt[tpv2[findfirst(x->x>v, tpv2)]]
186     end
187
188     ## initialization
189     neg_thr = sum(A[:, 3] .* (A[:, 3] .< 0))
190     d = fill(Inf, calV); d[s] = 0
191     pred = Vector{Int64}(undef, calV); pred[s] = 0
192     if use_dequeue
193         LIST = Deque{Int64}(); push!(LIST, s)
194         LIST_hist = Queue{Int64}(); enqueue!(LIST_hist, s)
195     else
196         LIST = Queue{Int64}(); enqueue!(LIST, s)
197     end
198     flag_nc = false
199
200     ## algorithm
201     n_ex = 0 # the number of node examinations
202     time::Float64 = 0
203     if use_dequeue
204         time = @elapsed while !isempty(LIST)
205             i = popfirst!(LIST)
206             for ind in pnt[i]:1:(pnt[i+1]-1)
207                 j = A[ind, 2]
208                 c_ij = A[ind, 3]
209                 n_ex += 1
210                 # check optimality condition
211                 if d[j] > d[i] + c_ij
212                     # update label
213                     d[j] = d[i] + c_ij
214                     # negative cycle detection
215                     if d[j] < neg_thr
216                         flag_nc = true
217                         break
218                     end
219                     pred[j] = i
220                     if !(j in LIST)
221                         if j in LIST_hist
222                             pushfirst!(LIST, j)
223                         else
224                             push!(LIST, j)
225                             enqueue!(LIST_hist, j)
226                         end
227                     end
228                 end
229             end
230         end
231     else
232         time = @elapsed while !isempty(LIST)
233             i = dequeue!(LIST)

```

```

234     for ind in pnt[i]:1:(pnt[i+1]-1)
235         j = A[ind, 2]
236         c_ij = A[ind, 3]
237         n_ex += 1
238         # check optimality condition
239         if d[j] > d[i] + c_ij
240             # update label
241             d[j] = d[i] + c_ij
242             # negative cycle detection
243             if d[j] < neg_thr
244                 flag_nc = true
245                 break
246             end
247             pred[j] = i
248             if !(j in LIST)
249                 enqueue!(LIST, j)
250             end
251         end
252     end
253 end
254 end
255 end
256
257 return d, pred, n_ex, time, flag_nc
258 end

```

The inputs for the algorithms are as follows:

- A: A matrix of integer variables representing the graph. Each row corresponds to an edge, where the first column is the tail node index, the second is the head node index, and the third is the edge's length.
- s: An integer variable representing the source node index from which the paths originate.
- use_dequeue: A boolean variable indicating the choice of data structure. If set to true, the algorithm employs a dequeue; a queue is used if false.

In the preprocessing phase (lines 177 to 186), several supporting variables are initialized before the algorithm execution. Here, the number of vertices (or the cardinality of the vertex set), denoted as calV , and the number of edges (or the cardinality of the edge set), denoted as calE , are defined. Additionally, in line with the star forward representation described by Ahuja et al. (1993, p. 34), the pointer vector, pnt , is also defined.

During the initialization phase (lines 188 to 198), all node labels are set to infinity except for the source node, which is labeled 0. This step also involves preparing containers for predecessors, nodes pending examination, and those already examined at least once when employing a dequeue strategy. A flag for detecting negative cycles is also introduced, set to true if a negative cycle is identified and false otherwise.

In the core algorithm part (lines 200 to 255), the implementation closely follows the pseudocode and explanations related to dequeue utilization as provided by Ahuja et al. (1993, pp. 141-143). The fundamental principle involves sequentially checking the optimality condition and updating any label that violates this condition. The counts of optimality condition checks (or node examinations) are tracked in the variables n_ex . The algorithm terminates if a negative cycle is detected or if every edge meets the optimality condition.

Upon completion, the function returns the labels and predecessors for each node, the total number of node examinations and label corrections, the computation time, and whether a negative cycle is present or not.

3 Experiment design and instances generation

The experiments were conducted across diverse network structures, categorized into 19 distinct groups based on three primary characteristics: the number of nodes, the number of arcs, and the presence (and the quantity) of edges with negative lengths. These groupings are detailed in Table 1. The first nine groups encompass networks where all edges possess non-negative lengths, while groups 10 through 19 include networks with varying edges with negative lengths. Despite the diversity of the network structures, all the nodes of each instance are reachable from the source node (with index 1).

Table 1: Network Structure of Each Instances Group

Inst. Group No.	# Nodes	# Edges	# Neg. Edges
1	10	18	0
2	10	45	0
3	10	72	0
4	50	490	0
5	50	1225	0
6	50	1960	0
7	100	1980	0
8	100	4950	0
9	100	7920	0
10	100	1980	20
11	100	1980	40
12	100	1980	99
13	100	1980	158
14	100	1980	198
15	100	4950	50
16	100	4950	99
17	100	4950	248
18	100	4950	396
19	100	4950	495

Groups 1 to 9 focus on networks that only include edges with non-negative lengths. These groups are subdivided based on the number of nodes, with 10 nodes in Groups 1 to 3, 50 in Groups 4 to 6, and 100 in Groups 7 to 9. Within each node-count category, the groups are ordered by an ascending number of edges. The edge lengths for these networks range from 0 to 50.

Groups 10 to 19 shift the focus towards networks with 100 nodes each but are unique in their incorporation of edges with negative lengths. These groups are segmented into two tiers based on edge density: Groups 10 to 14 are less dense, whereas Groups 15 to 19 feature more edges. Within both tiers, the number of edges with negative lengths

increases. For these groups, the range of edge lengths extends from -5 to 50.

To statistically analyze the algorithms' performance, each group comprised 1000 randomly generated instances.

4 Experiment results

The experiment is performed in my personal computer located in my laboratory, N7-2 3332. The specifications of the computer, language, and used packages are in Table 2

Table 2: Experimental Environment

Resource	Specification
CPU	13th Gen Intel(R) Core(TM) i9-13900K
RAM	128GB
OS	Windows 11
Language	Julia 1.10.0
Used packages	Random, Base, Statistics
	StatsBase (v0.33.21)
	DataStructures (v0.18.18)
	DelimitedFiles (v1.9.1)
	ProgressBars (v1.5.1)

The experiment results are in Table 3. The averages of time and the number of node examinations of each group with different data structures are provided. The values in the relative ratio are the dequeue's values relative to the queue's, given in percentage. These relative ratios provide a direct comparison: a negative sign indicates that the queue's performance metrics are higher (worse) than the dequeue's, and a positive sign indicates the opposite, where the dequeue's metrics are higher (worse) than those of the queue.

Analyzing the results reveals two observations. First, employing a queue consistently yields better computation times in every scenario involving exclusively non-negative arc lengths. However, for Groups 1, 2, and 4, which feature relatively simpler network structures, utilizing a queue surpasses the queue in terms of the number of basic operations.

Furthermore, in scenarios involving networks with negative arc lengths, the greater the number of negative arcs, the better the performance of using dequeue.

5 Conclusions

To summarize the findings:

- Using dequeue performs better on a moderate-size or sparse network with non-negative arc lengths.
- Using dequeue performs better on a network containing many arcs with negative lengths.

Table 3: Experiments Results, Average of N=1000 Instances for Each Group (Rel. Ratio = (dequeue's-queue's)/queue's)

Group No.	# Node Exam.			Time (Sec)		
	Dequeue	Queue	Rel. Ratio	Dequeue	Queue	Rel. Ratio
1	18.6	18.8	-1.1	5.9e-7	4.8e-7	24.7
2	53.1	54.5	-2.6	8.2e-7	6.7e-7	23.1
3	96.0	94.7	1.4	9.1e-7	7.9e-7	14.9
4	804.0	814.1	-1.2	8.0e-6	6.3e-6	26.7
5	2576.4	2399.2	7.4	1.3e-5	1.0e-5	26.0
6	4918.1	4182.2	17.6	1.6e-5	1.3e-5	27.1
7	4401.9	4141.4	6.3	3.1e-5	2.3e-5	36.4
8	13542.8	11989.8	13.0	5.4e-5	4.0e-5	35.5
9	24551.8	21030.7	16.7	7.1e-5	5.2e-5	36.4
10	48694.1	61792.6	-21.2	2.4e-4	2.4e-4	0.1
11	71142.3	80907.6	-12.1	3.8e-4	3.6e-4	5.9
12	85284.5	111829.7	-23.7	5.8e-4	5.8e-4	0.2
13	88553.8	142770.2	-38.0	7.2e-4	8.1e-4	-10.1
14	86043.9	160159.9	-46.3	7.8e-4	9.6e-4	-19.0
15	134393.7	185703.3	-27.6	5.1e-4	5.2e-4	-1.7
16	174228.8	268568.6	-35.1	7.8e-4	8.1e-4	-4.3
17	214607.4	419130.9	-48.8	1.4e-3	1.5e-3	-10.3
18	243736.8	483383.8	-49.6	1.9e-3	2.0e-3	-6.8
19	230921.0	499116.2	-53.7	2.0e-3	2.3e-3	-10.6

These observations lend cautious support to the statements made by Ahuja et al. (1993, p.143), particularly concerning networks with negative arc lengths.

References

Ahuja, R. K., Magnanti, T. L., and Orlin, J. B. (1993). *Network Flows: Theory, Algorithms, and Applications*. Pearson.